

CSW Proof-of-Value Findings — customer1

Generated 2026-05-01T18:30:00Z from live evidence in `https://csw.customer1.example`. Use this as the talk-track for the readout meeting; it is intentionally denser than the CISO brief.

What we looked at

In a single read of the tenant we collected:

Signal	Count
Workloads inventoried	1,880
Workloads reporting via agent	1,206 (64.1%)
Scopes in hierarchy	18
Policies authored	84
Policies in enforce mode	22
Absolute policies (highest precedence)	4
Critical CVEs on workloads	19
High CVEs on workloads	142
Enforcement alerts in 24h	187

The same evidence bundle is replayable; nothing here was hand-curated.

Posture at a glance

Score: 29 / 100 (Grade F)

- **Visibility — 16 pts.** 1206/1880 workloads reporting (64.1% agent coverage).
 - **Segmentation — 6 pts.** 84 policies, 22 in enforce mode (26%), 4 absolute (5%).
 - **Hygiene — 3 pts.** 19 critical, 142 high CVEs detected on workloads.
 - **Operations — 1 pts.** 187 enforcement alerts in last 24h.
 - **Compliance — 3 pts.** 7/25 controls marked 'met' across 4 frameworks.
-

What you're already doing well

Evidence collected for 4 frameworks

Re-using a single set of CSW evidence across multiple frameworks is the headline value proposition - this tenant is realising it.

What we'd recommend addressing

62 policies remain in monitor mode

Category: segmentation · Estimated posture uplift if addressed: +11 pts

Monitor-mode policies generate evidence but do not block lateral movement. Promoting reviewed policies to enforce directly improves the segmentation score component.

674 workloads without an agent

Category: visibility · Estimated posture uplift if addressed: +9 pts

Workloads outside agent coverage are invisible to flow analytics, ADM, and policy enforcement. Closing this gap unlocks visibility-driven controls in every framework.

19 critical CVEs on workloads

Category: hygiene · Estimated posture uplift if addressed: +9 pts

CSW correlates package CVEs to running processes. Each critical CVE on an internet-reachable workload is a candidate for escalation to patching or compensating policy.

Few absolute policies in place

Category: segmentation · Estimated posture uplift if addressed: +6 pts

Absolute policies are the right tool for crown-jewel isolation (PHI, cardholder data, regulated systems). Their absence indicates the segmentation strategy is still permissive-by-default.

187 enforcement alerts in 24h

Category: operations · Estimated posture uplift if addressed: +5 pts

Sustained alert volume indicates either policy gaps (legitimate flows blocked) or active attack pressure. Both warrant SOC engagement and policy review.

HIPAA: 5 controls not marked met

Category: compliance · Estimated posture uplift if addressed: +2 pts

5 controls in the HIPAA catalog still show partial or pending evidence. Address these before the next audit cycle.

Suggested 30-day plan

The plan is sequenced so each week's work feeds the next; weeks 1-2 are foundation moves, week 3 is the highest-leverage policy promotion, and week 4 closes the loop with a re-measurement.

Wk	Action	Owner	Addresses
1	Close agent-coverage gap. Identify the 674 workloads without an agent. Triage into: (a) deploy now, (b) excluded by policy, (c) decommission. Track in a single ticket; aim for 95%+ coverage.	Platform Engineering	674 workloads without an agent
2	Promote vetted policies to enforce mode. Run the 30-day flow review on monitor-mode policies for top 3 scopes. Promote those with no false-positive flow violations to enforce.	Security Operations	62 policies remain in monitor mode
3	Author absolute policies for crown jewels. Identify the top 3 crown-jewel scopes (databases of regulated data, payment systems, PHI stores). Draft absolute policies that take precedence over inherited rules.	Security Architecture	Few absolute policies in place
4	Address top	Vulnerability	19 critical CVEs

Wk	Action	Owner	Addresses
	critical CVEs. Pull the CSW vulnerability dashboard, sort by internet exposure + severity, dispatch top 10 to patching or compensating-policy workstreams.	Management	on workloads
5	Tune enforcement alert noise. Categorise the past 24h of alerts (legitimate block, missing exception, real attack). Build exception rules for the first two; route the third into incident response.	Security Operations	187 enforcement alerts in 24h
6	Refresh evidence and report. Re-run csw-insights against the tenant. Compare posture score to baseline; share delta with steering committee.	Compliance / GRC	—

How CSW shows up across the frameworks you care about

For each framework below, the runbooks and reports in CSW-Compliance-Mapping walk through the specific controls this evidence supports. The numbers here are the *current* state of controls evaluated in this tenant.

Framework	Controls evaluated	Met	Coverage
HIPAA	7	2	29%

Framework	Controls evaluated	Met	Coverage
PCI-DSS	6	1	17%
NIST-800-53	8	3	38%
SOC2	4	1	25%

Demo-able artefacts

The same evidence bundle that produced this report can also drive:

- **ADM (Application Dependency Mapping) walk-through** — visualise the highest-blast-radius workloads we identified.
- **Policy what-if** — replay 24h of flow data against a candidate policy to show enforcement impact before promoting.
- **Vulnerability dashboard** — sort the critical CVEs above by internet exposure to pick the top 10 patching candidates.
- **Auditor-ready evidence bundle** — the signed JSON used to render this report drops directly into the framework runbooks.

What we'd like to do next

1. Walk through the top 3 risks above with your platform/security teams.
2. Pick one risk to address this sprint; re-run `csw-insights` after.
3. Use the resulting before/after delta as the proof point for the wider rollout business case.

Materials, scripts, and templates are in CSW-Tenant-Insights. The underlying compliance narrative is in CSW-Compliance-Mapping. The evidence collection runs from CSW-API-Compliance-Automation.